

# Fundamentals of Database Systems

## COMPSCI 351

Instructor: Sebastian Link

The University of Auckland

# SQL

## As a Query Language

# What are Query Languages?

Query languages (QLs) give us

access and retrieval of data from a database

Relational model supports simple, powerful QLs:

- an operational language that allows DBs to speed up query evaluation
- a declarative language with a strong formal foundation based on logic
- an industry standard language for specifying queries in real DBMSs

Query Languages are not programming languages

- QLs not expected to be “Turing complete”
- QLs not intended to be used for complex calculations
- QLs support easy, efficient access to large data sets

# Formal Relational Query Languages

Two mathematical QLs form the basis for

- industry standards (e.g. SQL), and
- implementation

## Relational Algebra

more operational, very useful for representing execution plans

## Relational Calculus

- lets users describe what they want, rather than how to compute it
- i.e., non-operational, but declarative
- simple to write queries, and have them translated into SQL

Both languages can essentially express the same queries

## Definition (Query Language (QL))

A **query language** is

- a formal language  $\mathcal{L}$  such that
- each query  $Q \in \mathcal{L}$  is associated with
  - an *input-schema*  $in(Q)$ ,
  - an *output-schema*  $out(Q)$ , and
  - a *query mapping*  $q(Q) : inst(in(Q)) \rightarrow inst(out(Q))$   
↳ i.e., taking databases over  $in(Q)$  to databases over  $out(Q)$

# Comparison of Query Languages

Two queries  $Q_1$  and  $Q_2$  are **equivalent** (denoted  $Q_1 \equiv Q_2$ ), if

- they have the same input-schema  $in(Q_1) = in(Q_2)$ ,
- the same output-schema  $out(Q_1) = out(Q_2)$ , and
- the same query mapping  $q(Q_1) = q(Q_2)$ , i.e.,
  - $\hookrightarrow$  for each database  $db \in inst(in(Q_1))$  we have  $q(Q_1)(db) = q(Q_2)(db)$

Let  $\mathcal{L}$  and  $\mathcal{L}'$  be QLs

- $\mathcal{L}'$  **dominates**  $\mathcal{L}$  (notation:  $\mathcal{L} \sqsubseteq \mathcal{L}'$ ), if
  - $\hookrightarrow$  for each query  $Q \in \mathcal{L}$  there exists a query  $Q' \in \mathcal{L}'$  with  $Q' \equiv Q$
- $\mathcal{L}'$  and  $\mathcal{L}$  are **equivalent** (notation:  $\mathcal{L} \equiv \mathcal{L}'$ ), if
  - $\hookrightarrow \mathcal{L} \sqsubseteq \mathcal{L}'$  and  $\mathcal{L}' \sqsubseteq \mathcal{L}$  both hold

## Example: The Same Query in Different Languages

Consider our database schema from before

- MOVIE(title, year, country, run\_time, genre), DIRECTOR(id, title, year)
- PERSON(id, first\_name, last\_name, year\_born), ACTOR(id, title, year, role)

Query in English

What are the movies directed by 'Akira Kurosawa'?

SQL

```
SELECT  m.title, m.year
FROM    MOVIE m, DIRECTOR d, PERSON p
WHERE   m.title=d.title AND m.year=d.year AND
        d.id=p.id AND p.first_name='Akira' AND p.last_name='Kurosawa';
```

# Example: The Same Query in Different Languages

Consider our database schema from before

- MOVIE(title, year, country, run\_time, genre), DIRECTOR(id, title, year)
- PERSON(id, first\_name, last\_name, year\_born), ACTOR(id, title, year, role)

Query in English

What are the movies directed by 'Akira Kurosawa'?

Relational algebra

$$\pi_{\text{title, year}}(\sigma_{\text{last\_name}='Kurosawa'}(\sigma_{\text{first\_name}='Akira'}(\text{MOVIE} \bowtie \text{DIRECTOR} \bowtie \text{PERSON})))$$



# Example: The Same Query in Different Languages

Consider our database schema from before

- MOVIE(title, year, country, run\_time, genre), DIRECTOR(id, title, year)
- PERSON(id, first\_name, last\_name, year\_born), ACTOR(id, title, year, role)

Query in English

What are the movies directed by 'Akira Kurosawa'?

Relational calculus

$$\{(x_{\text{title}}, x_{\text{year}}) \mid \exists x_{\text{country}}, x_{\text{run\_time}}, x_{\text{genre}}, x_{\text{id}}, x_{\text{year\_born}} (\text{MOVIE}(x_{\text{title}}, x_{\text{year}}, x_{\text{country}}, x_{\text{run\_time}}, x_{\text{genre}}) \wedge \text{DIRECTOR}(x_{\text{id}}, x_{\text{title}}, x_{\text{year}}) \wedge \text{PERSON}(x_{\text{id}}, \text{'Akira'}, \text{'Kurosawa'}, x_{\text{year\_born}}))\}$$

# SELECT... FROM... WHERE...

The basic form for all queries is given by the SELECT statement

```
SELECT <attribute-list> FROM <table-list> WHERE <condition> ;
```

## Comments

- \* in the attribute-list means all attributes
- Attribute names may be qualified with the table name
  - required, if attribute-names are not unique
- Attribute and table names can be renamed: <name> AS <short name>
  - in this case the short name is required for qualification

# Duplicates and Order

SQL is based on multisets (instead of sets)  
`SELECT DISTINCT` removes duplicates

The result may be ordered adding an `ORDER BY` clause  
`ORDER BY` <attribute-list>

Add `ASC` or `DESC` for ascending or descending order

# Example

## English Language Query

- List all people with their id, first name, last name and the year they were born
- Provided their id is between 221 and 300 or greater than 999.
- Order the list first by last name, then by first name.

## Corresponding SQL Query

```
SELECT      *  
FROM        PERSON  
WHERE       (id >= 221 AND id <= 300) OR id >= 1000  
ORDER BY    last_name, first_name;
```

# Example

## English Language Query

List the first and last name of all movie directors of non-US movies together with the titles and production years of these movies.

## Corresponding SQL Query

```
SELECT  p.first_name, p.last_name, m.title, m.year
FROM    MOVIE m, PERSON p, DIRECTOR d
WHERE   d.title = m.title AND m.country <> 'USA' AND
        d.year = m.year AND d.id = p.id;
```

## Condition in WHERE-clause (in queries, deletes, updates)

- Can be a simple comparison:  $\langle \text{expression} \rangle \langle \text{operator} \rangle \langle \text{expression} \rangle$
- The comparison operators can be  $=$ ,  $<>$ ,  $<=$ ,  $>=$ ,  $>$  or  $<$
- Expressions built from attributes, constant values and operators defined for the corresponding domains (see manuals for details)

# More Complex WHERE-Conditions

## Other conditions

- value lists:
  - ↳ `<attribute> BETWEEN <value> AND <value>`, or
  - ↳ `<attribute> IN <value-list>`
- `IS NULL` and `IS NOT NULL` test for null markers
- `<attribute> LIKE <string>` allows pattern matching
  - ↳ uses `_` for a single symbol and `%` for any substring

WHERE-clause may be **complex condition** using AND, OR and NOT

# Example

## English Language Query

List all movies together with their directors, provided the last name of the director begins with 'H' and ends with an 'o' followed by two more letters.

## Corresponding SQL Query

```
SELECT  m.title, p.first_name, p.last_name
FROM    MOVIE m, PERSON p, DIRECTOR d
WHERE   m.title = d.title AND m.year = d.year AND
        d.id = p.id AND p.last_name LIKE 'H%o__';
```



# Example

## English Language Query

List all movies produced between 1978 and 1995, ordered by year

## Corresponding SQL Query

```
SELECT    title, year
FROM      MOVIE
WHERE     year BETWEEN 1978 and 1995
ORDER BY  year;
```

# Aggregate functions

Aggregate functions applied to attributes or attribute lists result in single output value

## Some functions

- COUNT returns the total number of input values
- AVG returns the average of input values, given some number domain
- MIN returns the minimum value of the inputs
- MAX returns the maximum value of the inputs
- SUM returns the sum of the input values, given some number domain

Several functions can be used in a query, but not in combination with simple attributes

The result is significantly different, if **DISTINCT** is used

# Example

## English language query

How many actors are stored in the database?

## Corresponding SQL query

```
SELECT  COUNT (DISTINCT a.id) AS number_of_actors
FROM    ACTOR a
```

# Example

## English language query

List minimum, average and maximum run-time of all non-German movies

## Corresponding SQL query

```
SELECT  MIN (run_time) AS min_run_time,  
        AVG (run_time) AS avg_run_time,  
        MAX (run_time) AS max_run_time  
FROM    MOVIE  
WHERE   country <> 'Germany';
```

# Grouping

Apply aggregate function to selected groups of tuples, for instance the number of actors for each movie

`GROUP BY <attribute-list>`

for building groups of tuples for each value combination in the attribute-list

The attributes in the attribute-list following the `GROUP BY` must be exactly the simple attributes in the list following the `SELECT` clause

# Example

## English language query

List each movie produced in 1999 together with its number of actors

## Corresponding SQL query

```
SELECT      m.title, m.year COUNT(DISTINCT a.id) AS number_of_actors
FROM        MOVIE m, ACTOR a
WHERE       a.title = m.title AND a.year = 1999 AND m.year = 1999
GROUP BY    m.title, m.year;
```

# The HAVING clause

The clause HAVING  $\langle \text{condition} \rangle$  evaluates a condition on the groups

The HAVING condition must apply to the group, not to individual tuples in the group

# Example

List for all countries except Australia the number of movies, if there are at least two.

```
SELECT      m.country, COUNT(m.title) AS number_of_movies
FROM        MOVIE m
WHERE       m.country <> 'Australia'
GROUP BY    m.country
HAVING      COUNT(*) >= 2;
```



# Sub-queries

## Sub-queries

are just queries that are used within WHERE-conditions

## IN ⟨sub-query⟩

tests if a tuple occurs in the result of a sub-query

## EXISTS ⟨sub-query⟩

tests whether sub-query results in non-empty relation

## Sub-queries continued

UNIQUE ⟨sub-query⟩

tests if result of sub-query contains no duplicates

ALL, SOME or ANY (SOME=ANY)

before a sub-query makes sub-queries usable in comparison formulae

Negation

In all these cases the condition involving the sub-query can be negated using a preceding NOT

# Example

## English Language Query

List for all countries except Spain the movies produced in that country, provided there are at least two of them. Order the result first ascending by country, then descending by the production year of the movie.

## Corresponding SQL Query

```
SELECT    m1.country, m1.year, m1.title
FROM      MOVIE m1
WHERE     m1.country <> 'Spain' AND m1.country IN (
            SELECT m2.country
            FROM MOVIE m2
            GROUP BY m2.country
            HAVING COUNT(*) >= 2)
ORDER BY  m1.country ASC, m1.year DESC;
```

# Example

## English Language Query

List the movies with exactly one actor, and return the first and last name of this actor.

## Corresponding SQL Query

```
SELECT  m.title, m.year, p.first_name, p.last_name
FROM    MOVIE m, ACTOR a1, PERSON p
WHERE   m.title = a1.title AND m.year = a1.year AND a1.id = p.id AND NOT EXISTS(
        SELECT *
        FROM Actor a2
        WHERE a2.id <> a1.id AND a2.title = a1.title AND a2.year = a1.year);
```

# Example

## English Language Query

List the NZ movies together with their run time where the run time is longer than any Australian movie.

## Corresponding SQL Query

```
SELECT  m1.title, m1.year, m1.run_time
FROM    MOVIE m1
WHERE   m1.country = 'New Zealand' AND m1.run_time > ANY (
        SELECT m2.run_time
        FROM MOVIE m2
        WHERE m2.country='Australia');
```

# Example

## English Language Query

List the actors who played in most movies

## Corresponding SQL Query

```
SELECT  p.first_name, p.last_name, p.id, a2.num_of_movies
FROM    PERSON p1, (
        SELECT a1.id, COUNT(DISTINCT a1.title, a1.year) AS num_of_movies
        FROM ACTOR a1
        GROUP BY a1.id ) a2
WHERE   p.id = a2.id AND a2.num_of_movies = (
        SELECT MAX(a4.num_of_movies)
        FROM (SELECT a3.id, COUNT(DISTINCT a3.title, a3.year) as num_of_movies
              FROM ACTOR a3
              GROUP BY a3.id ) a4);
```

# Limitations of SQL

FLIGHTS

Origin	Destination
Singapore	Tokyo
Singapore	Frankfurt
Singapore	Auckland
Tokyo	Joburg
	⋮

Reachability Query:

Find pairs of cities ( $O, D$ ) such that one can fly from  $O$  to  $D$  with at most one stop

```
SELECT  F1.Origin, F2.Destination
FROM    FLIGHTS F1, FLIGHTS F2
WHERE   F1.Destination = F2.Origin
UNION
SELECT  *
FROM    FLIGHTS;
```

## Another Reachability Query

Find pairs of cities ( $O, D$ ) such that one can fly from  $O$  to  $D$  with at most two stops

```
SELECT  F1.Origin, F3.Destination
FROM    FLIGHTS F1, FLIGHTS F2, FLIGHTS F3
WHERE   F1.Destination = F2.Origin AND F2.Destination = F3.Origin
UNION
SELECT  F1.Origin, F2.Destination
FROM    FLIGHTS F1, FLIGHTS F2
WHERE   F1.Destination = F2.Origin
UNION
SELECT  *
FROM    FLIGHTS;
```



# Reachability Queries in General

For any fixed  $k$ , we can write in SQL

Find  $(O, D)$  so one can fly from  $O$  to  $D$  with at most  $k$  stops.

What about general reachability?

Find pairs of cities  $(O, D)$  such that one can fly from  $O$  to  $D$ .

SQL'99 solution

New construct to express reachability

SQL'92

is not able to express this query

This feature is not implemented in all products

# A Different Point of View

## Rewrite reachability queries as rules

$$\text{REACH}(x, y) : \neg \text{FLIGHTS}(x, y)$$
$$\text{REACH}(x, y) : \neg \text{FLIGHTS}(x, z), \text{REACH}(z, y)$$

Latter rule is *recursive*: REACH refers to itself

## Evaluation

- Step 0:  $\text{REACH}_0$  is initialised as the empty set
- Step  $i + 1$ :

$$\text{REACH}_{i+1}(x, y) : \neg \text{FLIGHTS}(x, y)$$
$$\text{REACH}_{i+1}(x, y) : \neg \text{FLIGHTS}(x, z), \text{REACH}_i(z, y)$$

- Stop condition:  
if  $\text{REACH}_{i+1} = \text{REACH}_i$ , then return it as answer of query

# Example of Evaluation

## FLIGHTS

Origin	Destination
Singapore	Tokyo
Tokyo	Joburg
Joburg	Rio

## Evaluation steps

Stop condition:  $REACH_4 = REACH_3$

## REACH

$\{(Singapore, Tokyo), (Tokyo, Joburg), (Joburg, Rio), (Singapore, Joburg), (Tokyo, Rio), (Singapore, Rio)\}$

## SQL'99 syntax mimics that of recursive rules

```
WITH RECURSIVE REACH(Origin, Destination) AS
(
    SELECT *
    FROM FLIGHTS
UNION
    SELECT F.Origin, R.Destination
    FROM FLIGHTS F, REACH R
    WHERE F.Destination = R.Origin
)
SELECT *
FROM REACH
```

# Recursion in SQL'99: Syntactic restrictions

A different way to describe reachability as a recursive rule-based query:

$$\begin{aligned}\text{REACH}(x, y) &: -\text{FLIGHTS}(x, y) \\ \text{REACH}(x, y) &: -\text{REACH}(x, z), \text{REACH}(z, y)\end{aligned}$$

This translates into an SQL'99 query

```
WITH RECURSIVE REACH(Origin, Destination) AS (  
  SELECT * FROM FLIGHTS  
  UNION  
  SELECT R1.Origin, R2.Destination  
  FROM REACH R1, REACH R2  
  WHERE R1.Destination = R2.Origin)  
SELECT * FROM REACH
```

Many implementations support  
*linear* recursion only

recursively defined relation  
mentioned just once in FROM line

## Another Example of Recursion in SQL'99

Query: Find cities reachable from Singapore

Assume `FLIGHTS` contains the further attribute `Aircraft`

```
WITH CITIES AS SELECT Origin, Destination FROM FLIGHTS
  RECURSIVE REACH(Origin, Destination) AS
  (
    SELECT * FROM CITIES
  UNION
    SELECT C.Origin, R.Destination
    FROM CITIES C, REACH R
    WHERE C.Destination = R.Origin
  )
SELECT R.Destination
FROM REACH R
WHERE R.Origin = "Singapore"
```

# Problems with Negation in Recursion

Consider rule such as  $R(X) : \neg S(X), \neg R(X)$  in SQL

```
WITH RECURSIVE R(A) AS
  (SELECT S.A
   FROM S
   WHERE S.A NOT IN (SELECT R.A FROM R))
SELECT * FROM R
```

For  $S = \{1, 2\}$  evaluation leads to...

after step 0:  $R_0 = \emptyset$ ;

after step 1:  $R_1 = \{1, 2\}$ ;

...

after step  $2n$ :  $R_{2n} = \emptyset$ ;

after step  $2n + 1$ :  $R_{2n+1} = \{1, 2\}$ ;

Problem: the evaluation does not terminate

negation (NOT IN) causes this problem

# Summary for SQL

- SQL is the industry standard for defining and querying data
- DDL of SQL allows us to declare tables, views, and constraints
- DML of SQL allows us to specify and execute queries and updates
- SQL permits duplicate and partial information
  - ↳ uses list semantics rather than set semantics of relational model
- SQL has several additional features, including
  - ↳ Indices
  - ↳ Access rights
  - ↳ Aggregate functions
  - ↳ Transitive closures
  - ↳ Embedded SQL
  - ↳ Dynamic SQL
  - ↳ Transactions